

tf-idf_function

March 15, 2025

```
[1]: # Libraries
%matplotlib inline
import pandas as pd
from collections import Counter
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.decomposition import LatentDirichletAllocation
from sklearn.decomposition import NMF
from sklearn.decomposition import TruncatedSVD
from sklearn.decomposition import PCA
from tqdm import tqdm
tqdm.pandas()
pd.options.display.max_colwidth = 150 ###
import numpy as np
import re
import sys
import os
import matplotlib.pyplot as plt
# Add the directory containing visualization_utils.py to path
sys.path.append("/Users/debr/English-Homer/")
import visualization_utils as viz
import seaborn as sns
sns.set_style("whitegrid")
# palette astroblue    orange    genoa    carrot    tawny    neptune    ↵
    ↪SELAGO    mako    black
color = ['#003D59', ↵
    ↪'#FD6626', '#177070', '#FB871D', '#641B5E', '#86C3BC', '#F5E1FD', '#414A4F', 'k']
danB_plotstyle = {'figure.figsize': (12, 7),
                  'axes.labelsize': 'large', # fontsize for x and y labels (was ↵
    ↪large)

                  'axes.titlesize': 'large', # fontsize for title
                  'axes.titleweight': 'bold', # font type for title
                  'xtick.labelsize': 'large', # fontsize for x
                  'ytick.labelsize': 'small', # fontsize for y ticks
                  'grid.color': 'k', # grid color
                  'grid.linestyle': ':', # grid line style
                  'grid.linewidth': 0.2, # grid line width
```

```

        'font.family': 'Times New Roman', # font family
        'grid.alpha': 0.5, # transparency of grid
        'figure.dpi': 300, # figure display resolution
        'savefig.bbox': 'tight', # tight bounding box
        'savefig.pad_inches': 0.4, # padding to use when saving
        'axes.titlepad': 15, # title padding
        'axes.labelpad': 8, # label padding
        'legend.borderpad': .6, # legend border padding
        'axes.prop_cycle': plt.cycler(
            color=color) # color cycle for plot lines
    }

# adjust matplotlib defaults
plt.rcParams.update(danB_plotstyle)

```

```

[2]: # Load CSVs
filepath_Wilson = "/Users/debr/odysseys_en/Odyssey_dfs/Odyssey_Wilson_eda_END.
    ↪CSV"
filepath_Green = "/Users/debr/odysseys_en/Odyssey_dfs/Odyssey_Green_eda_END.csv"

df_W = pd.read_csv(filepath_Wilson)
df_G = pd.read_csv(filepath_Green)

# Add translation label
df_W["translation"] = "Wilson"
df_G["translation"] = "Green"

# merging "book_num" with "translation" to create a unique identifier
df_W["book_id"] = df_W["book_num"].astype(str) + "_W"
df_W = df_W.drop(columns=["book_num"])
df_G["book_id"] = df_G["book_num"].astype(str) + "_G"
df_G = df_G.drop(columns=["book_num"])

# Keep only necessary columns: book number & tokens
df_W = df_W[["book_id", "tokens"]]
df_G = df_G[["book_id", "tokens"]]

# Combine both into one DataFrame
df = pd.concat([df_W, df_G], ignore_index=True)

# Ensure tokens are stored as lists (if stored as strings, convert them)
df["tokens"] = df["tokens"].apply(lambda x: eval(x) if isinstance(x, str) else_
    ↪x)

```

```

[3]: import pandas as pd
import numpy as np
from collections import Counter

```

```

def calculate_tfidf(df):
    """
    Calculate TF-IDF scores for a DataFrame with book_id and tokens columns.

    Parameters:
    -----
    df : pandas DataFrame
        A DataFrame with 'book_id' and 'tokens' columns.
        The 'tokens' column should contain lists of tokens (as strings or
        ↪ actual lists).

    Returns:
    -----
    pandas DataFrame
        The original DataFrame with additional columns:
        - term_freq: Dictionary of term frequencies for each token
        - term_counts: Dictionary of raw counts for each token
        - idf: Dictionary of IDF scores for each token
        - tf_idf: Dictionary of TF-IDF scores for each token
    """
    # Create a copy of the DataFrame to avoid modifying the original
    result_df = df.copy()

    # Function to compute term frequency and term counts
    def term_freq_by_doc(list_of_tokens):
        # Handle both string representation of list and actual list
        if isinstance(list_of_tokens, str):
            token_list = eval(list_of_tokens) # Convert string representation
            ↪ to list
        else:
            token_list = list_of_tokens # Use as is if already a list

        # Count occurrences of each term
        term_counts = Counter(token_list)

        # Total number of terms in the document
        total_terms = len(token_list)

        # Compute TF: term frequency for each token
        term_freq = {term: count / total_terms for term, count in term_counts.
            ↪ items()}

        return term_freq, term_counts

    # Apply function to compute TF for each book

```

```

    result_df["term_freq"], result_df["term_counts"] = zip(*result_df["tokens"].
    ↪apply(term_freq_by_doc))

    # Get total number of documents (books)
    N = len(result_df)

    # Count how many documents contain each term
    doc_containing_term = Counter()
    for term_counts in result_df["term_freq"]:
        doc_containing_term.update(term_counts.keys()) # Count unique terms in
    ↪each document

    # Compute IDF for each term
    idf_scores = {term: np.log(N / (1 + doc_count)) for term, doc_count in
    ↪doc_containing_term.items()} # Adding 1 to avoid division by zero

    # Add IDF column to df
    result_df["idf"] = result_df["term_freq"].apply(lambda term_freq: {term:
    ↪idf_scores[term] for term in term_freq})

    # Compute TF-IDF by multiplying TF and IDF for each term in each document
    result_df["tf_idf"] = result_df.apply(lambda row: {term:
    ↪row["term_freq"][term] * row["idf"][term] for term in row["term_freq"]},
    ↪axis=1)

    return result_df

# Example usage:
df_tfidf_W = calculate_tfidf(df_W)
df_tfidf_G = calculate_tfidf(df_G)

```

[4]: # Top terms for each book and heatmap plot

```

def extract_top_terms(df, n=50):
    """
    Extract the top N most important terms from the tf_idf column

    Parameters:
    -----
    df : pandas DataFrame
        DataFrame with 'book_id' and 'tf_idf' columns
    n : int
        Number of top terms to extract (default: 50)

    Returns:
    -----
    tuple
    """

```

```

        (top_terms_per_book, top_terms_overall)
        - top_terms_per_book: DataFrame with top terms for each book
        - top_terms_overall: DataFrame with top terms across all books
    """
    # Extract top terms per book
    top_terms_per_book = {}

    for _, row in df.iterrows():
        book_id = row['book_id']
        tf_idf_dict = row['tf_idf']

        # Sort terms by tf-idf score (descending) and take top N
        sorted_terms = sorted(tf_idf_dict.items(), key=lambda x: x[1],
        ↪reverse=True)[:n]
        top_terms_per_book[book_id] = {term: score for term, score in
        ↪sorted_terms}

    # Convert to DataFrame for easier analysis
    top_terms_df = pd.DataFrame.from_dict(top_terms_per_book, orient='index')

    # Extract top terms overall
    all_terms = {}
    for tf_idf_dict in df['tf_idf']:
        for term, score in tf_idf_dict.items():
            if term in all_terms:
                all_terms[term] += score
            else:
                all_terms[term] = score

    # Sort terms by total tf-idf score (descending) and take top N
    top_terms_overall = sorted(all_terms.items(), key=lambda x: x[1],
    ↪reverse=True)[:n]

    # Convert to DataFrame
    top_terms_overall_df = pd.DataFrame(top_terms_overall, columns=['term',
    ↪'total_score'])

    return top_terms_df, top_terms_overall_df

def create_tfidf_heatmap(df, top_n=50):
    """
    Create a heatmap of the top N terms across all books

    Parameters:
    -----
    df : pandas DataFrame
        DataFrame with 'book_id' and 'tf_idf' columns

```

```

top_n : int
    Number of top terms to include in the heatmap (default: 50)
"""
# Extract top terms overall
_, top_terms = extract_top_terms(df, n=top_n)
top_terms_list = top_terms['term'].tolist()

# Create a matrix of book_id x top_terms
heatmap_data = []
book_ids = []

for _, row in df.iterrows():
    book_id = row['book_id']
    book_ids.append(book_id)

    tf_idf_dict = row['tf_idf']

    # Extract scores for top terms
    scores = [tf_idf_dict.get(term, 0) for term in top_terms_list]
    heatmap_data.append(scores)


# Convert to numpy array
heatmap_array = np.array(heatmap_data).T

# Create heatmap
plt.figure(figsize=(14, 16))
sns.heatmap(heatmap_array, cmap='Blues', xticklabels=book_ids,
yticklabels=top_terms_list)
plt.title(f'Top {top_n} Terms by TF-IDF Score ({df["book_id"].iloc[0][2:
])}')
plt.xlabel('Terms')
plt.ylabel('Books')
plt.xticks(rotation=0)
plt.tight_layout()
plt.savefig(f"/Users/debr/English-Homer/MVP_Green-Wilson/plots/
tfidf_heatmap({df['book_id'].iloc[0][2:]).png")
plt.show()


return heatmap_array

# Example usage
top_terms_per_book_W, top_terms_overall_W = extract_top_terms(df_tfidf_W)
heatmap_array = create_tfidf_heatmap(df_tfidf_W)

```



```
[5]: top_terms_per_book_G, top_terms_overall_G = extract_top_terms(df_tfidf_G)
heatmap_array = create_tfidf_heatmap(df_tfidf_G)
```





```
[6]: # Getting the 30 top terms for each translation
tt_W = top_terms_overall_W['term'][:30]
tt_G = top_terms_overall_G['term'][:30]
# New df with the top terms from both for comparison
top_terms_overall = pd.DataFrame({'Wilson': tt_W, 'Green': tt_G})
top_terms_overall
```



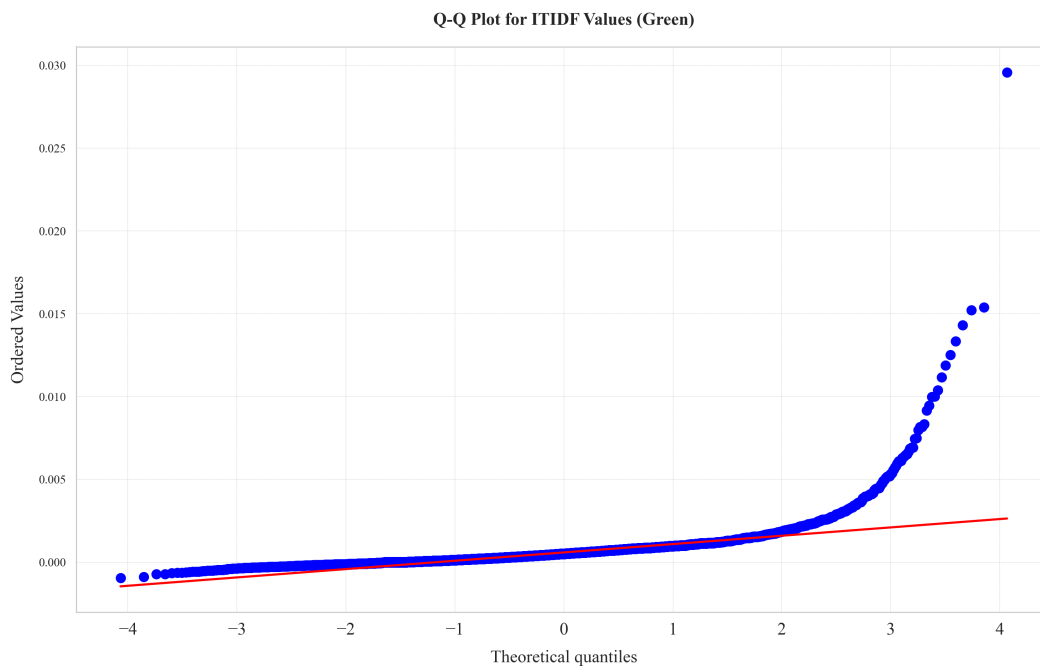
```
[6]:      Wilson      Green
0      bow      alkinoos
1      alcinous      bow
2      swineherd      swineherd
3      circe      telemachos
4      telemachus      phaiakians
5      eumaeus      kirke
6      nestor      nestor
7      suitors      '
8      '      eumaios
9      menelaus      suitors
10     irus      antinoos
11     antinous      raft
12     raft      menelaos
13     scylla      nausikaa
14     string      iros
15     old      cave
16     aegisthus      girl
17     melanthius      goddess
18     horses      aigisthos
19     poseidon      wagon
20     nausicaa      string
21     arete      melanthios
22     cave      poseidon
23     phaeacia      skylle
24     mentor      arete
25     calypso      stranger
26     achilles      atreus
27     agamemnon      penelope
28     mules      eurymachos
29     river      bend
```


```
[7]: itidf_W = df_tfidf_W["tf_idf"].apply(lambda x: list(x.values())).sum()
      itidf_G = df_tfidf_G["tf_idf"].apply(lambda x: list(x.values())).sum()
```

```
[8]: import scipy.stats as stats
      # Q-Q plot for ITIDF values
      stats.probplot(itidf_W, dist="norm", plot=plt)
      plt.title("Q-Q Plot for ITIDF Values (Wilson)")
      plt.savefig(f"/Users/debr/English-Homer/MVP_Green-Wilson/plots/Q-Q_Plot_Wilson.
      ↪png")
      plt.show()
      sns.histplot(itidf_W, bins=30, kde=True) # KDE adds a smoothed curve
      plt.title("ITIDF Distribution (Wilson)")
      plt.savefig(f"/Users/debr/English-Homer/MVP_Green-Wilson/plots/
      ↪ITIDF_Distribution_Wilson.png")
      plt.show()
```



```
[9]: import scipy.stats as stats
# Q-Q plot for ITIDF values
stats.probplot(itidf_G, dist="norm", plot=plt)
plt.title("Q-Q Plot for ITIDF Values (Green)")
plt.savefig(f"/Users/debr/English-Homer/MVP_Green-Wilson/plots/Q-Q_Plot_Green.
↳png")
plt.show()
sns.histplot(itidf_G, bins=30, kde=True) # KDE adds a smoothed curve
plt.title("ITIDF Distribution (Green)")
plt.savefig(f"/Users/debr/English-Homer/MVP_Green-Wilson/plots/
↳ITIDF_Distribution_Green.png")
plt.show()
```





```
[10]: from scipy.stats import shapiro
# Shapiro test for normality
def shapiro_test(data):
    """
    Perform the Shapiro-Wilk test for normality from scipy.stats and print the
    results.

    Parameters:
    -----
    data : list or array
        The data to be tested for normality.

    Returns: shapiro statistic and p-value

    """
    stat, p = shapiro(data)
    print(f"Shapiro-Wilk Test: p-value = {p}")

    if p > 0.05:
        print("Data appears to be normally distributed (fail to reject H0).")
    else:
        print("Data is not normally distributed (reject H0).")
    return stat, p
```

```
# Shapiro-Wilk test for normality
print("Wilson")
shapiro_test(itidf_W)
print("\n")
print("Green")
shapiro_test(itidf_G)
```

Wilson

Shapiro-Wilk Test: p-value = 1.8664746605407613e-105

Data is not normally distributed (reject H0).

Green

Shapiro-Wilk Test: p-value = 1.5165217224910905e-113

Data is not normally distributed (reject H0).

```
[10]: (np.float64(0.6884581528049329), np.float64(1.5165217224910905e-113))
```

```
[11]: # Mannwhitney U test for comparing ITIDF values
from scipy.stats import mannwhitneyu # Import the actual function

def mannwhitneyu_test(x, y, alternative='two-sided'):
    """
    Perform the Mann-Whitney U test for comparing two independent samples.
    """
    stat, p = mannwhitneyu(x, y, alternative=alternative) # Use SciPy's
    ↪function
    print(f"Mann-Whitney U test statistic: {stat}, p-value: {p}")

    if p < 0.05:
        print("Reject H: The distributions of the translations are
        ↪significantly different.")
    else:
        print("Fail to reject H: No significant difference between the
        ↪translations.")

# Run the test
mannwhitneyu_test(itidf_W, itidf_G, alternative='two-sided')
```

Mann-Whitney U test statistic: 398708946.5, p-value: 4.071216635611771e-145

Reject H: The distributions of the translations are significantly different.